

Entity Linker

Touring Testers
May 2015

Hugo Sandelius, Tobias Norlund,
Simon Schweighofer, Fabian Stutz

Baselines

BASELINE 1

Development Set

F1 strict: **45%**

F1 relaxed: **49%**

Train Set

F1 strict: **40%**

F1 relaxed: **43%**

BASELINE 2

Development Set

F1 strict: **51%**

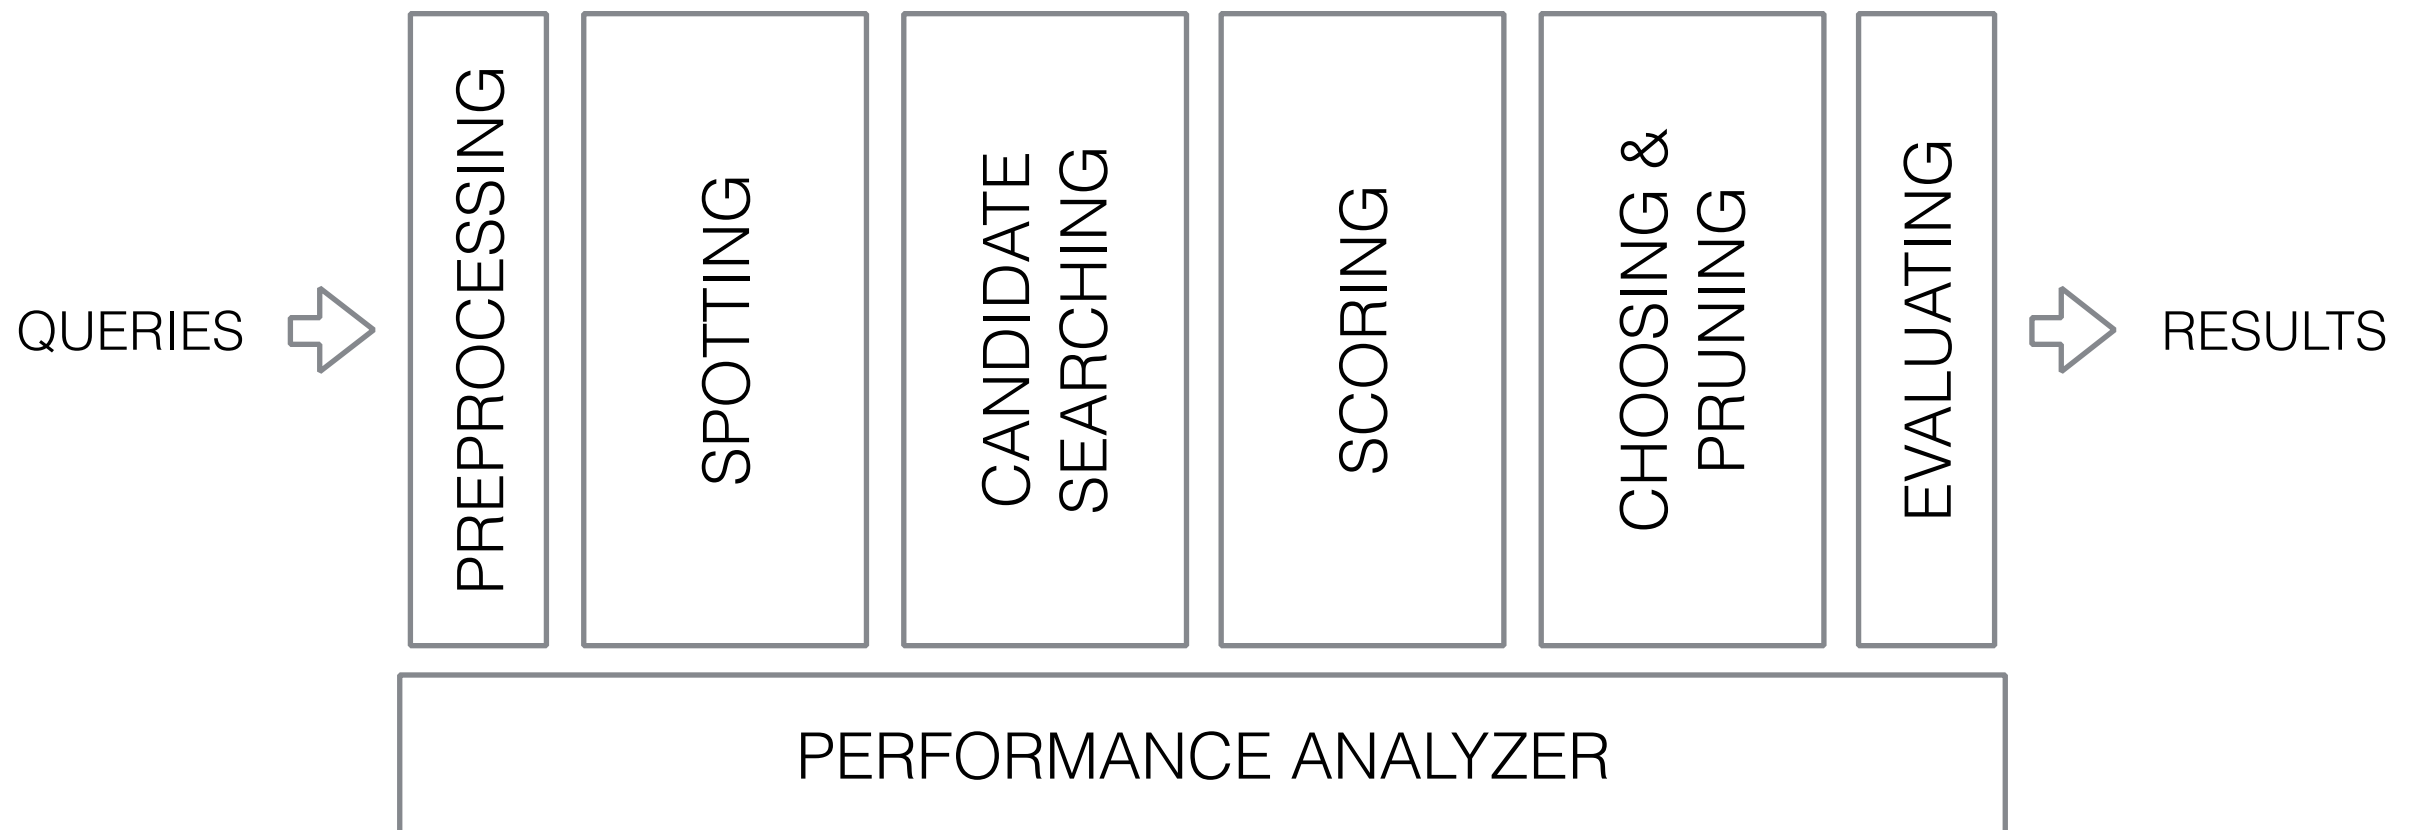
F1 relaxed: **59%**

Train Set

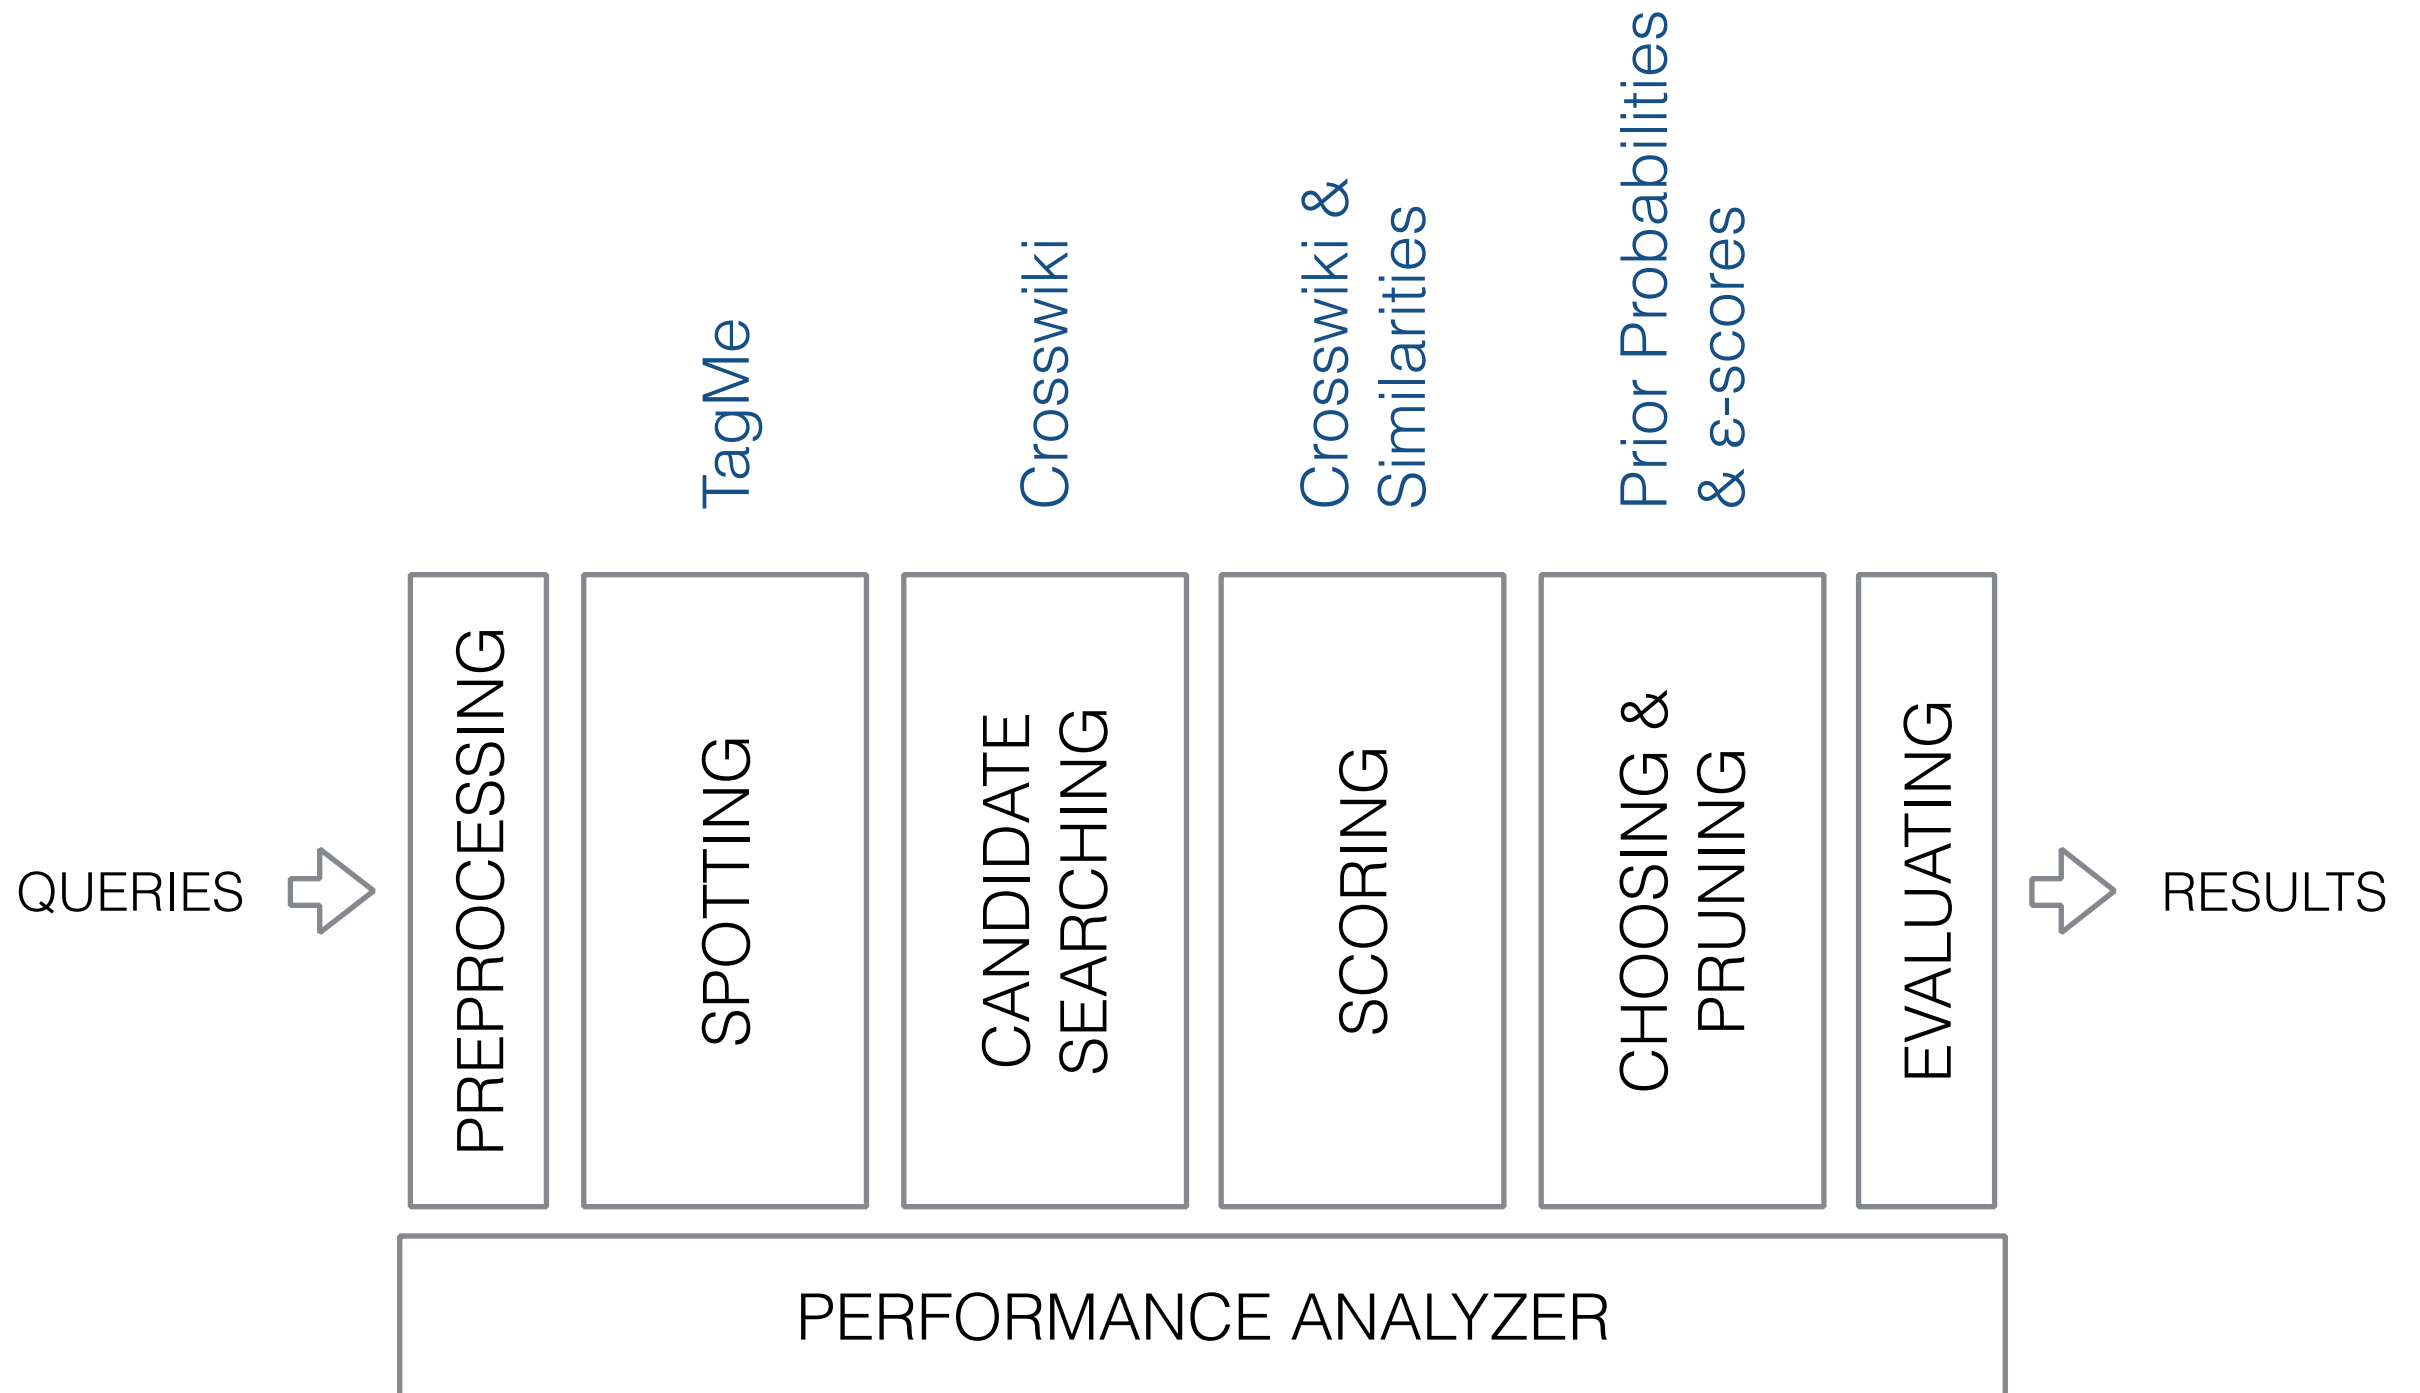
F1 strict: **51%**

F1 relaxed: **59%**

Architecture



Base Implementation



Performance Analyzer

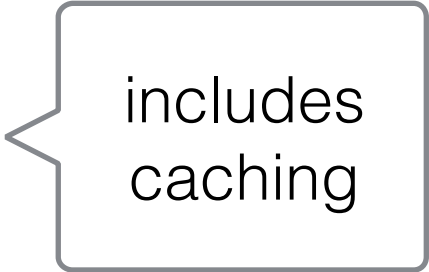
- **Compare Results with the Gold Standard**
 - Do we select the correct spots?
 - Do we look at the correct candidates?
 - Do we choose the correct candidates?
- **Find bottlenecks**

Performance Analyzer

- **Showed the following:**
 - Spotting is (quite) precise
 - Sometimes there are no candidates at all
 - „Gold“ candidate often not among candidates
 - The gold candidate often gets selected, if among candidates

Improvements

- **SPOTTING** (84% correct mentions)
 - small adjustments (still TagMe based)
- **CANDIDATE SEARCHING** (91% correct candidates)
 - ~~Crosswiki~~
 - Wikipedia search & suggestions
 - bing search for english wikipedia articles
- **SCORING**
 - Parameter adjustments
- **CHOOSING** (81% correct)
- **PRUNING**
 - Explored different approaches → not successful



includes
caching

Results

BASELINE 2

Development Set

F1 strict: **51%**

F1 relaxed: **59%**

Train Set

F1 strict: **51%**

F1 relaxed: **59%**

FINAL

Development Set

F1 strict: **58%**

F1 relaxed: **64%**

Train Set

F1 strict: **55%**

F1 relaxed: **59%**